

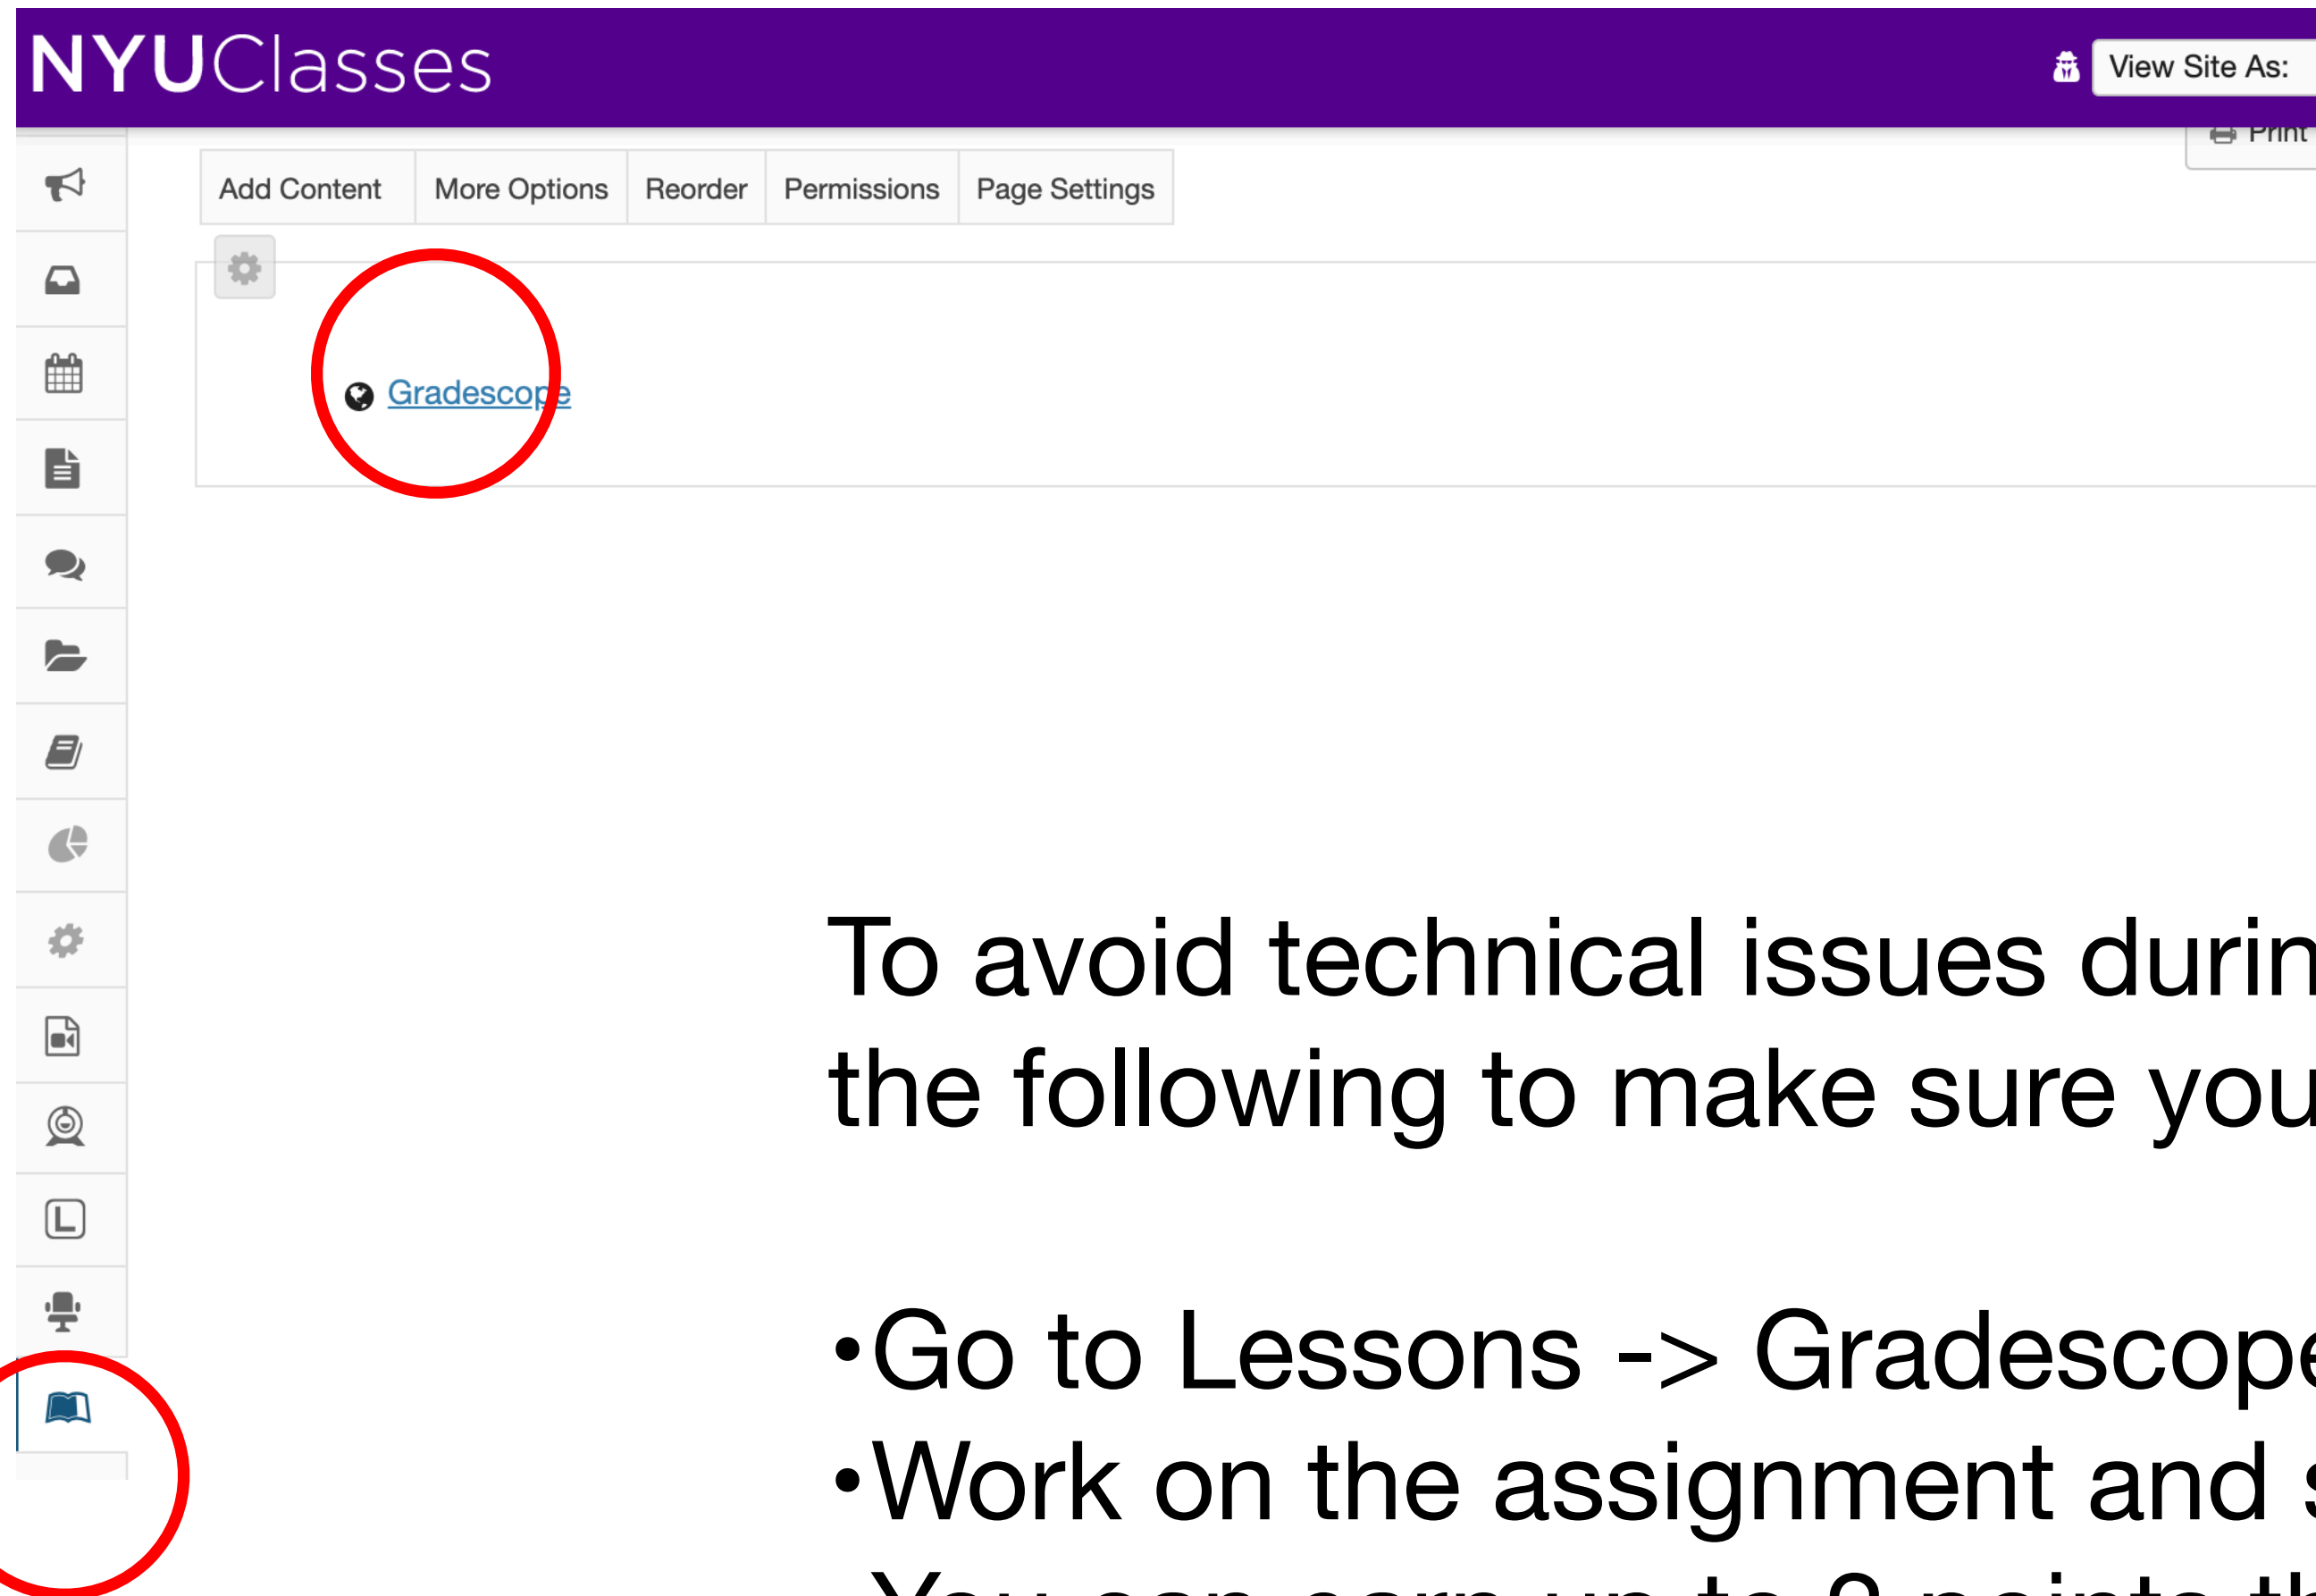
Today

- Last time
 - SVMs with kernels (reading: Bishop Sec 7.1)
 - Multiclass classification (reading: Bishop, Sec 7.1.3)
- Today
 - Optimization for SVM
 - Nearest neighbor classifiers (Bishop, Sec 2.5.2, Sec 6.3; Hastie, Sec 13)
 - Some learning theory (Ng, Lecture notes
<http://cs229.stanford.edu/notes/cs229-notes4.pdf>)
- Announcements
 - HW 2 due on Wed, Oct 14
 - Lab on Wed, Oct 14; blended (know your seat number)
 - Midterm exam: **Wed, Oct 21**, 11.00am - 12.15pm (New York time)
 - Q&A on Mon, Oct 19 in lecture
 - Midterm online via Gradescope (more details to come)

How to prepare for the midterm exam

- Lecture (+)
 - Good for getting the “big picture” and a general understanding
 - Work through the derivations and try to understand each step
- Homeworks (++)
 - Learn how to implement concepts discussed in lecture
 - Theory questions useful to test your understanding (expect similar style of questions on the exam)
- Textbooks have many more exercises (++++)
 - Bishop, Pattern Recognition and Machine Learning
 - Hastie, The elements of statistical learning (data mining, inference, and prediction)
 - Murphy, Machine learning
- Make the most out of study groups, office hours, Q&A session, etc.

Gradescope



To avoid technical issues during the midterm exam, try the following to make sure your setup is working properly:


- Go to Lessons -> Gradescope -> TestYourSetup assignment
- Work on the assignment and submit
- You can earn up to 2 points that are counted towards the midterm

The assignment is open from today until Wed before class.

Algorithm for solving SVM problems

board

Coordinate descent





SMO algorithm

board

Nonlinear classifiers

Material based on David Sontag's slides, Andrew Ng's notes, and Cho's notes

Recall: From linear to nonlinear decision boundary



[David Sontag]

Recall: Feature map

- Define a feature map $\phi : \mathbb{R}^n \rightarrow \mathbb{R}^k, k \in \mathbb{N} \cup \{\infty\}$
- Polynomial regression with degree 3

$$\phi(x) = \begin{bmatrix} x^3 \\ x^2 \\ x \\ 1 \end{bmatrix}$$

- Let's apply (linear) SVM to $\phi(x)$ instead of x
- Potential problem?
 - The dimension k potentially very (infinite) high
 - Goal: Avoid operating in \mathbb{R}^k

Example: XOR problem

0	0	1
0	1	0
1	0	0
1	1	1

Example: XOR problem (cont'd)

- This data set is ***not*** linearly separable





Figure: Cho

Feature mapping for XOR problem



Rotate with

$$R(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}$$

$$x^r = R(\pi/4)x$$

Still not linearly separable





Figure: Cho

Feature mapping for XOR problem (cont'd)

XOR (Original)



Take absolute value
of each rotated vector

$$\mathbf{x}^{ra} = |R(\pi/4)\mathbf{x}|$$
$$= \phi(\mathbf{x})$$



Now it is linearly
separable

XOR (Rotated+Abs)




Figure: Cho

Another transformation for XOR problem

- Select 4 “basis vectors”

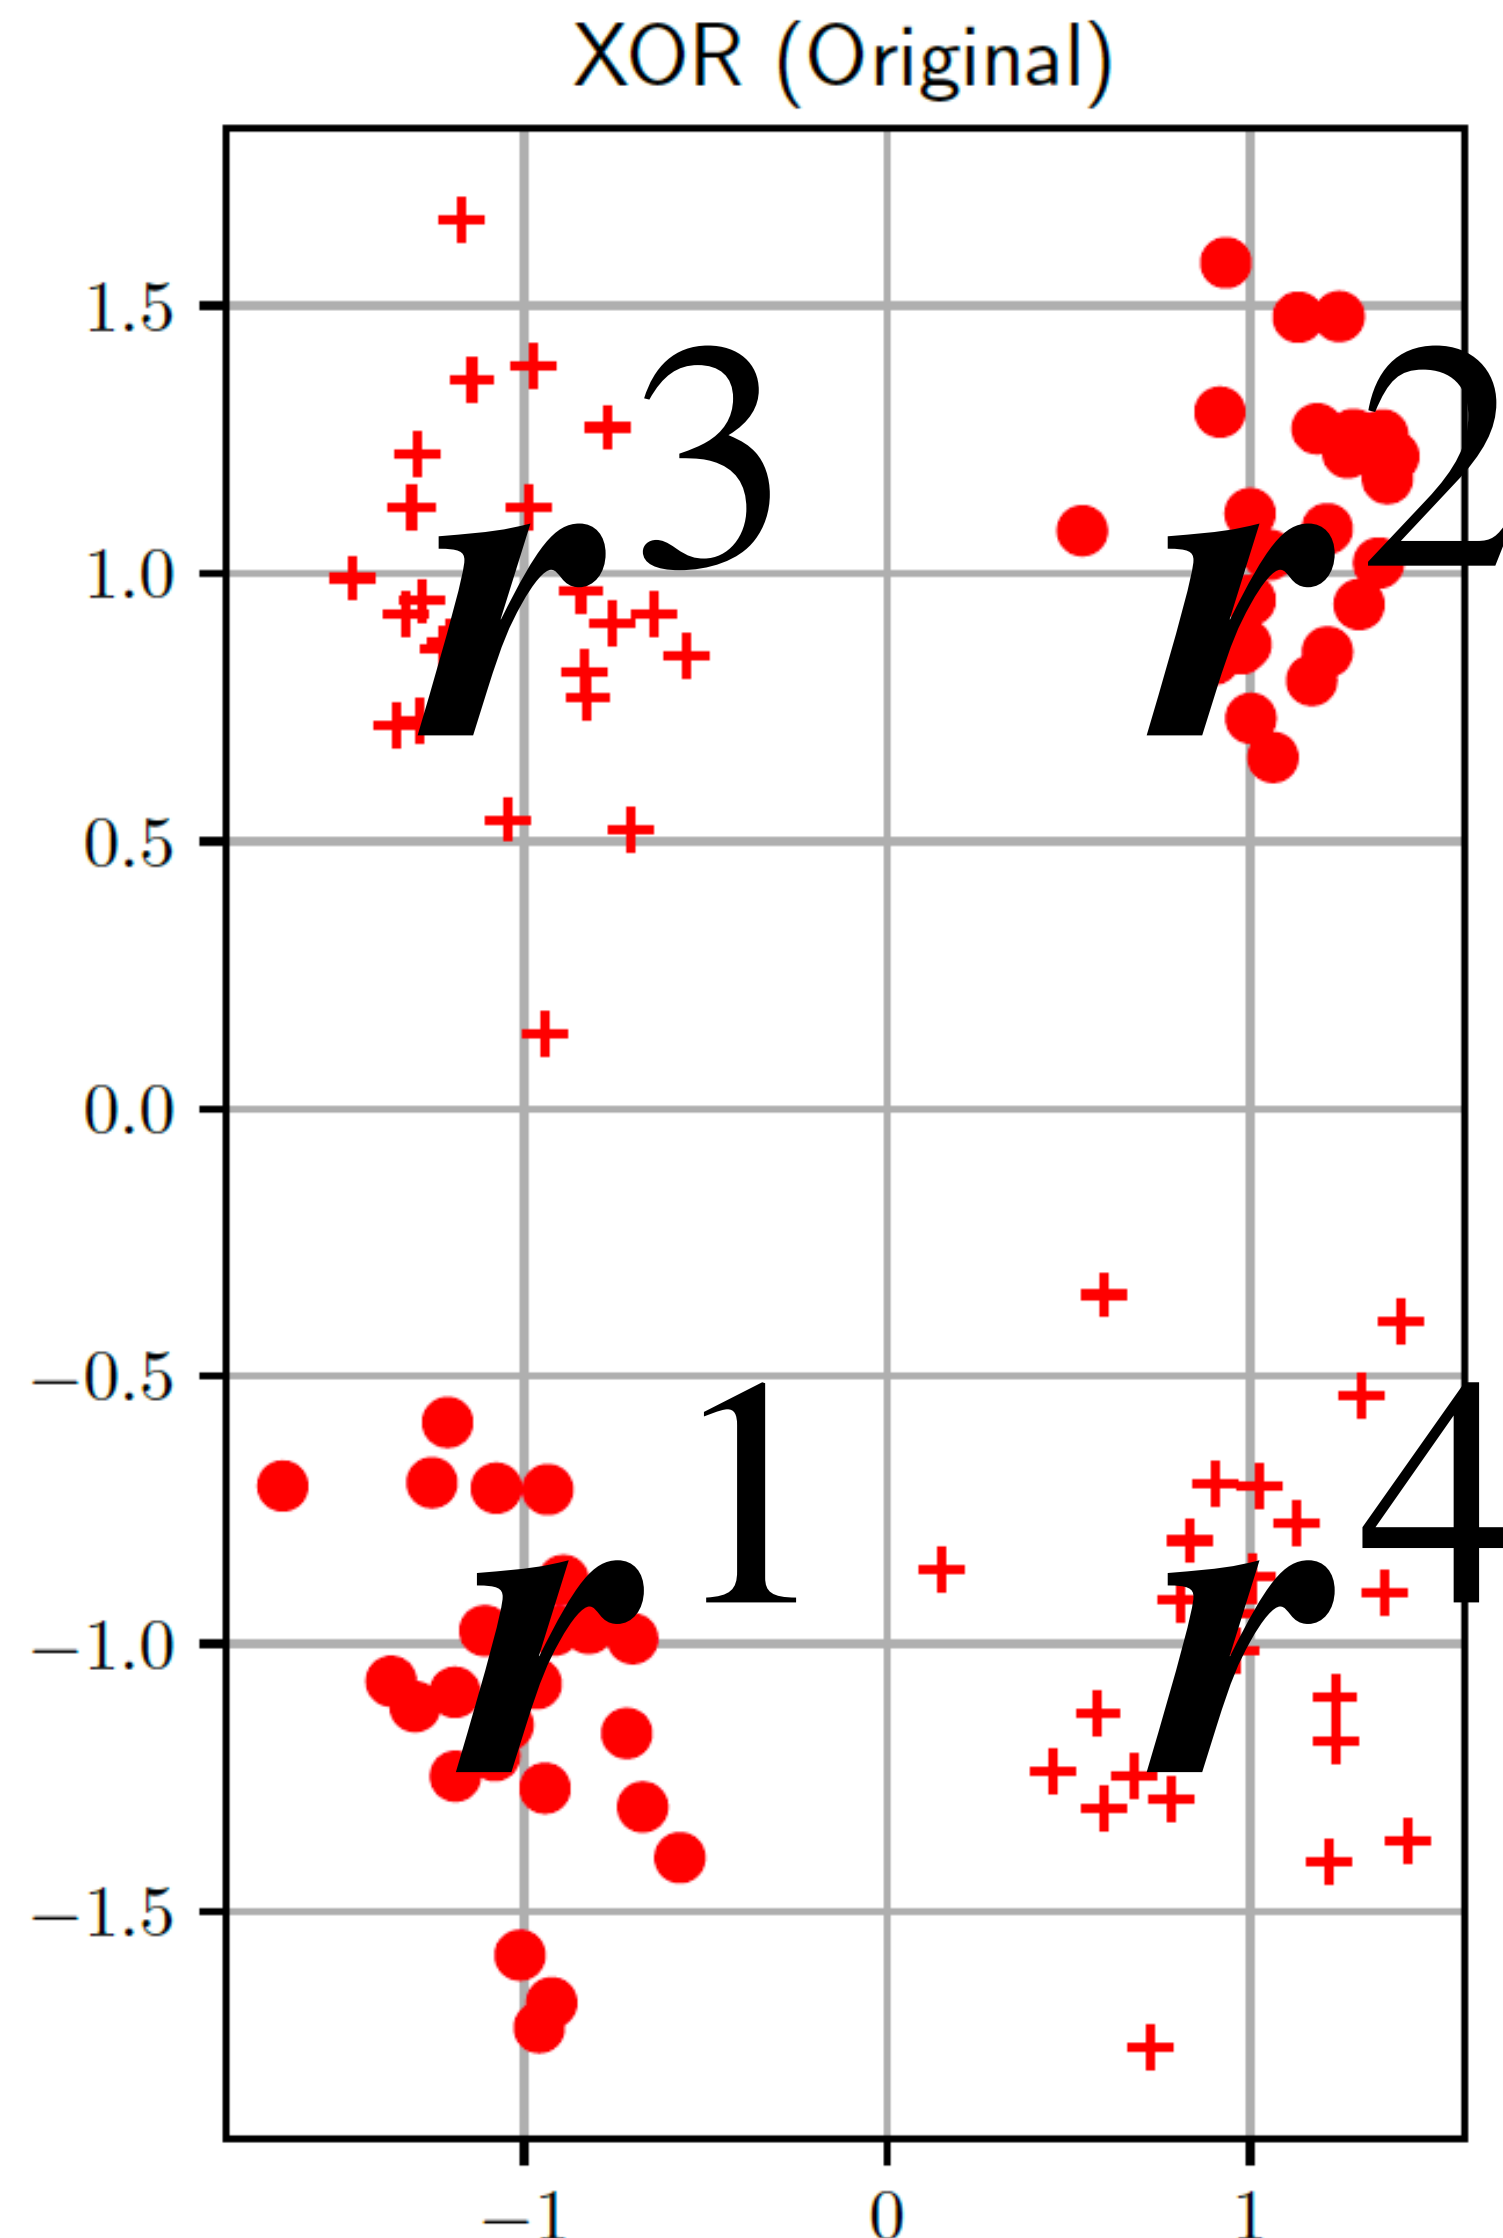
$$\mathbf{r}^1 = [-1, -1]^T, \mathbf{r}^2 = [1, 1]^T, \mathbf{r}^3 = [-1, 1]^T, \mathbf{r}^4 = [1, -1]^T$$

- Define the feature mapping

$$\phi(\mathbf{x}) = \begin{bmatrix} \phi_1(\mathbf{x}) \\ \phi_2(\mathbf{x}) \\ \phi_3(\mathbf{x}) \\ \phi_4(\mathbf{x}) \end{bmatrix} = \begin{bmatrix} \exp(-\|\mathbf{x} - \mathbf{r}^1\|^2) \\ \exp(-\|\mathbf{x} - \mathbf{r}^2\|^2) \\ \exp(-\|\mathbf{x} - \mathbf{r}^3\|^2) \\ \exp(-\|\mathbf{x} - \mathbf{r}^4\|^2) \end{bmatrix}$$

- The i -th component is inverse proportional to the distance between \mathbf{x} and vector \mathbf{r}^i

Closeness



- Close to r^1, r^2 should be classified as positive (circles)
 - Input vector in positive class are close to either r^1 or r^2
- Close to r^3, r^4 should be classified as negative (plus)
 - Input vector in positive class are close to either r^3 or r^4

Figure: Cho

Weight vector for linear classifier

- Class labels $y \in \{-1, 1\}$
- Parametrize with θ

$$h_{\theta}(x) = g(w^T x)$$

with

$$g(z) = \begin{cases} 1, & z \geq 0 \\ -1, & z < 0 \end{cases}$$

- Weight vector $\theta = [1, 1, -1, -1]^T$ perfectly solves XOR problem for transformed data $\phi(x)$
 - Proof: Board
- Note that $\theta = [y^{(1)}, y^{(2)}, y^{(3)}, y^{(4)}]$ is the vector containing the class label -1 or 1 of the basis vectors

Generalize this concept of constructing feature map

- Have K basis vectors with their labels

$$\{(\mathbf{r}^1, y^{(1)}), \dots, (\mathbf{r}^K, y^{(K)})\}$$

- Each input is transformed as

$$\phi(\mathbf{x}) = \begin{bmatrix} \exp(-\|\mathbf{x} - \mathbf{r}^1\|^2) \\ \vdots \\ \exp(-\|\mathbf{x} - \mathbf{r}^K\|^2) \end{bmatrix}$$

- With weight vector

$$\boldsymbol{\theta} = [y^{(1)}, \dots, y^{(K)}]^T$$

Nearest neighbor classifier

- Have N training points (= basis vectors) with their labels

$$\{(\mathbf{x}^1, y^{(1)}), \dots, (\mathbf{x}^N, y^{(N)})\}$$

- Each input is transformed as

$$\phi(\mathbf{x}) = \begin{bmatrix} \exp(-\|\mathbf{x} - \mathbf{x}^1\|^2) \\ \vdots \\ \exp(-\|\mathbf{x} - \mathbf{x}^N\|^2) \end{bmatrix}$$

- Set only the max component to 1: $\phi'_k(x) = \begin{cases} 1, & \phi_k(x) = \max \phi(x) \\ 0, & \text{else} \end{cases}$

- With weight vector

$$\boldsymbol{\theta} = [y^{(1)}, \dots, y^{(N)}]^T$$

- This is a nearest neighbor classifier $h_{\boldsymbol{\theta}}(\mathbf{x}) = g(\boldsymbol{\theta}^T \phi'(\mathbf{x}))$

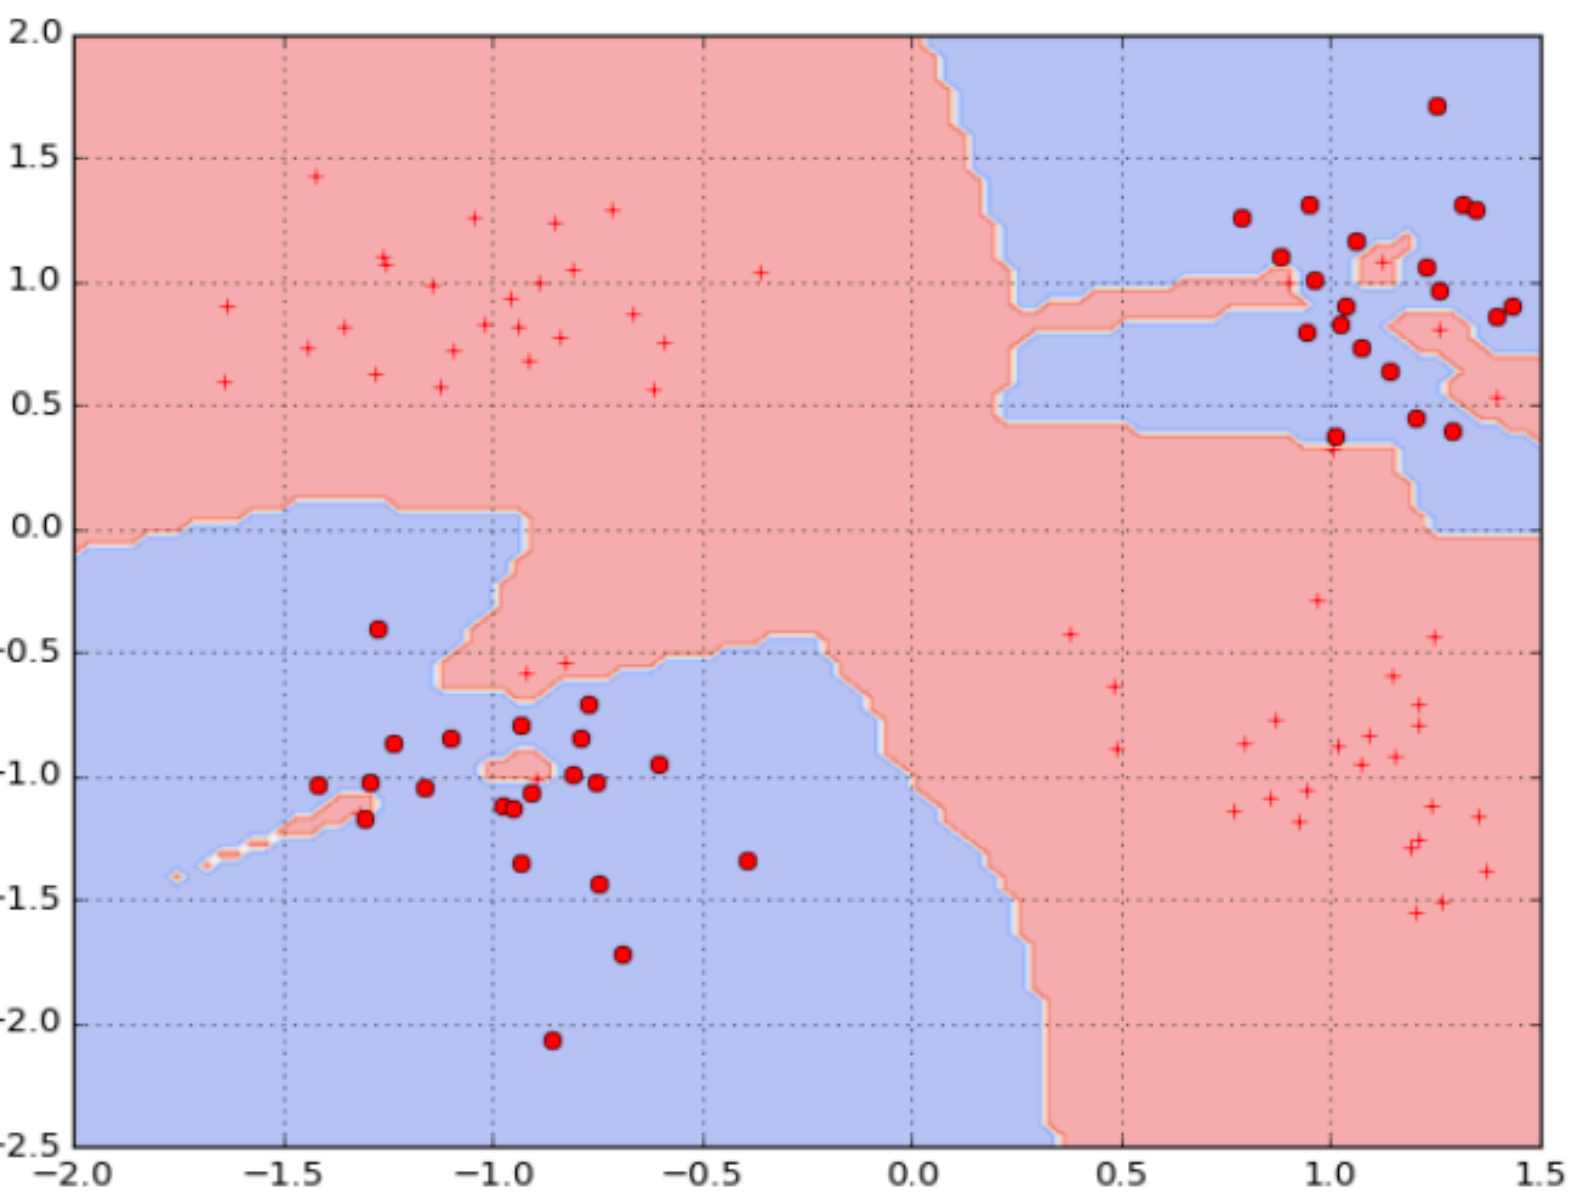
Different point of view on nearest neighbor

- Instead of feature map, search for nearest neighbor and return label of nearest neighbor
- Different distance than Euclidean $\| \cdot \|_2$ possible
- Rarely use “nearest neighbor”. Rather use k -nearest neighbor (KNN)
 - Search k nearest neighbors
 - Vote which label new point gets
- How can we interpret number of neighbors k ?


Different point of view on nearest neighbor

- Instead of feature map, search for nearest neighbor and return label of nearest neighbor
- Different distance than Euclidean $\| \cdot \|_2$ possible
- Rarely use “nearest neighbor”. Rather use k -nearest neighbor (KNN)
 - Search k nearest neighbors
 - Vote which label new point gets
- The number of neighbors serves as regularizer
 - If $k = 1$ then costs low because only single neighbor required
 - However, leads to overfitting because a single point determines label in neighborhood
 - Other extreme is $k = N$, then over-regularization


Parameter k controls regularization



(a) $k = 1$



(b) $k = 5$



(c) $k = 20$

As k increases, the decision boundary becomes smoother

Selecting basis vectors - random

- Weakness of KNN
 - Requires potentially large storage (because need to store all N training points)
 - Need to sweep through entire training set for each predict
- Idea: Select only a few representative basis vectors
 - Uniform-randomly select an index from $i \in \{1, \dots, N\}$
 - Set $B = B \cup \{i\}$
 - Set $D = D \cup \{x^{(i)}\}$
 - Repeat if $|B| < k$
- Works surprisingly well
 - Basis vectors are close to training data set
 - Basis vectors “evenly” distributed in training data set (uniform sampling)
 - Weight vector θ without optimization (cheap)

Selecting basis vectors - clustering

- Earlier in the XOR problem, we first performed clustering and then picked for each cluster a representative basis vector
 - Difficult to find clustering non-visually in $>3D$
 - Will discuss clustering in detail later
- Once clusters have been selected, need to set weight vector θ

- The k basis vectors r^1, \dots, r^k define ϕ

$$\phi(\mathbf{x}) = \begin{bmatrix} \exp(-\|\mathbf{x} - \mathbf{r}^1\|^2) \\ \vdots \\ \exp(-\|\mathbf{x} - \mathbf{r}^k\|^2) \end{bmatrix}$$

- **No labels** for basis vectors r^1, \dots, r^k given
 - Gives new training data set $D_\phi = \{(\phi(x^{(1)}), y^{(1)}), \dots, (\phi(x^{(N)}), y^{(N)})\}$
 - Apply **linear** classifier to D_ϕ (logistic regression, SVM, etc)
- Sometimes known under name “radial basis function network”

Learning the feature map (towards deep learning)

- So far, we said given are basis vectors that define feature map
- Then, we found the weights via linear classification
- In principle, we can parametrize the basis vectors as well and then optimize for the right basis vectors via gradient descent
- Further, we could use another ϕ than the one we used based on radial basis functions
 - In principle, any differentiable ϕ is feasible
 - We could think of one ϕ that does the feature map in one shot
 - Or we could compose many simpler ones

$$\phi = \phi_\ell(\cdots\phi_1(x)) = (\phi_\ell \circ \dots \circ \phi_1)(x)$$

- This is a topic of deep learning \Rightarrow see other courses